

Linear regression: Abalone growth and pH example

Luke Simon

2026-05-15

Table of contents

Set up	1
Abalone example	2
a. Questions and hypotheses	2
b. Cleaning	2
c. Exploratory data visualization	2
d. Abalone model	3
Model fitting	3
Diagnostics	4
Model summary	4
Generating predictions	5
Why generate model predictions?	6
Visualizing model predictions and data	7
Creating a table with model coefficients, 95% confidence intervals, and more . .	8

Set up

```
library(tidyverse) # general use
library(janitor) # cleaning data frames
library(here) # file/folder organization
library(readxl) # reading .xlsx files
library(ggeffects) # generating model predictions
library(flextable) # creating tables
```

```
# abalone data from Hamilton et al. 2022
abalone <- read_xlsx(here("data", "Abalone IMTA_growth and pH.xlsx"))
```

Abalone example

Data from Hamilton et al. 2022. “Integrated multi-trophic aquaculture mitigates the effects of ocean acidification: Seaweeds raise system pH and improve growth of juvenile abalone.”
<https://doi.org/10.1016/j.aquaculture.2022.738571>

a. Questions and hypotheses

Question: How does pH predict abalone growth (measured in change in shell surface area per day, $\text{mm}^2 \text{d}^{-1}$)?

H_0 : pH has no effect on abalone growth.

H_A : pH has a


b. Cleaning

```
# creating new data frame
abalone_clean <- abalone |>
  # standardizing the names of columns
  clean_names() |>
  # selecting columns of interest
  select(mean_p_h, change_in_area_mm2_d_1_25) |>
  # renaming columns to make more standard
  rename(mean_ph = mean_p_h,
         change_in_area = change_in_area_mm2_d_1_25)
```

Don't forget to look at your data! Use `View(abalone_clean)` in the Console.

c. Exploratory data visualization

```
# base layer: ggplot
ggplot(data = abalone_clean, # data from clean data frame
       aes(x = mean_ph, # x axis: pH
           y = change_in_area)) + # y-axis: change in area
# first layer: data points
geom_point(size = 4,
           stroke = 1,
           fill = "firebrick4",
           shape = 21)
```



Where does the missing value come from?

Missing value comes from one pair of tanks where pH sensor malfunctioned.

Don't forget to turn your warnings off either in the individual code chunk or in the YAML at the top of the file.


d. Abalone model

Model fitting

```
abalone_model <- lm(
  change_in_area ~ mean_ph, # formula: change in shell SA/day as a function of pH
  data = abalone_clean      # data frame: cleaned abalone data
)
```

Diagnostics

```
par(mfrow = c(2, 2)) # displays plots in 2x2 grid
plot(abalone_model)   # displays diagnostic plots for model
```



Model summary

```
# more information about the model
summary(abalone_model)
```

Call:

```
lm(formula = change_in_area ~ mean_ph, data = abalone_clean)
```

Residuals:

Min	1Q	Median	3Q	Max
-0.42714	-0.29282	0.08769	0.21258	0.43965

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	-18.5010	6.1601	-3.003	0.01488 *
mean_ph	2.9827	0.7596	3.926	0.00348 **

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 0.3063 on 9 degrees of freedom

(1 observation deleted due to missingness)

Multiple R-squared: 0.6314, Adjusted R-squared: 0.5905

F-statistic: 15.42 on 1 and 9 DF, p-value: 0.003477

What is the slope estimate (with standard error?)

β_1 : 2.98 ± 0.76

What is the intercept estimate (with standard error?)

β_0 : -18.50 ± 6.16

Generating predictions

```
# creating a new object called abalone_preds
abalone_preds <- ggpredict(
  abalone_model,      # model object
  terms = "mean_ph"   # predictor (in quotation marks)
)
```

```
# display the predictions
```

```
abalone_preds
```

Predicted values of change_in_area

mean_ph	Predicted	95% CI
7.94	5.19	4.83, 5.54

7.95		5.21		4.86, 5.55
8.06		5.53		5.30, 5.76
8.10		5.67		5.46, 5.88
8.10		5.67		5.46, 5.88
8.14		5.77		5.55, 5.98
8.15		5.81		5.59, 6.04
8.33		6.36		5.92, 6.80

Look at the column names using `colnames(abalone_preds)` in the Console.

What are the column names?

`x`, `predicted`, `std.error`, `conf.low`, `conf.high` (not worrying about `group` for now)

Look at the actual object (by clicking on it in Environment).

Look at the class of the object using `class(abalone_preds)` in the Console.

What is this type of object?

It is a `ggeffects` output and a regular data frame.

View `abalone_preds` like a data frame.

Why generate model predictions?

This dataset does not include any actual observations of abalone growth at $\text{pH} = 8$.

However, the model can tell us a prediction for what abalone growth could be if $\text{pH} = 8$ based on the existing data.

Figure out what abalone growth the model would predict at $\text{pH} = 8$.

```
# finding the model prediction at a specific value
ggpredict(
  abalone_model,      # model object
  terms = "mean_ph[8]" # predictor (in quotation marks) and predictor value in brackets
)
```

```
# Predicted values of change_in_area
```


mean_ph		Predicted		95% CI
8		5.36		5.08, 5.64

How would you interpret this?

At a pH of 8 the model predicts change in abalone shell surface area to be 5.36 [95% CI: 5.08,. 5.64] ($\text{mm}^2 \text{ d}^{-1}$).

Visualizing model predictions and data

```
# base layer: ggplot
# using clean data frame
ggplot(data = abalone_clean,
       aes(x = mean_ph,
           y = change_in_area)) +
# first layer: points representing abalones
geom_point(size = 4,
           stroke = 1,
           fill = "firebrick4",
           shape = 21) +
# second layer: ribbon representing confidence interval
# using predictions data frame
geom_ribbon(data = abalone_preds,
          aes(x = x,
              y = predicted,
              ymin = conf.low,
              ymax = conf.high),
          alpha = 0.1) +
# third layer: line representing model predictions
# using predictions data frame
geom_line(data = abalone_preds,
         aes(x = x,
             y = predicted)) +
# axis labels
labs(x = "Mean pH",
     y = expression("Change in shell area ("*mm^{2}*d^{-1}*")")) +
# theme
theme_minimal()
```



Creating a table with model coefficients, 95% confidence intervals, and more

Note: this function is from `flextable`.

```
as_flextable(abalone_model)
```

	Estimate	Standard Error	t value	Pr(> t)
(Intercept)	-18.501	6.160	-3.003	0.0149 *
mean_ph	2.983	0.760	3.926	0.0035 **

*Signif. codes: 0 '***' < 0.001 < '**' < 0.01 < '*' < 0.05*

Residual standard error: 0.3063 on 9 degrees of freedom

Multiple R-squared: 0.6314, Adjusted R-squared: 0.5905

F-statistic: 15.42 on 9 and 1 DF, p-value: 0.0035

END OF ABALONE EXAMPLE